

# Comparing Observation and Action Representations for Reinforcement Learning in $\mu$ RTS

Shengyi Huang and Santiago Ontañón

Drexel University

*{sh3397,so367}@drexel.edu*

October 7, 2019

# Overview

- 1 Introduction
  - Background
  - Problem Statement
  - Existing Work
- 2 Observation and Action Representations
  - Global Representation
  - Local Representation
- 3 Advantage Actor Critic
  - A2C
  - Specifics
- 4 Evaluation
  - Experimental Setup
  - Evaluation
- 5 Conclusions

# RTS Games




Figure: StarCraft II (PC, 2010)

<https://youtu.be/yaqez9Snt4E?t=13>

# RTS Games

- RTS Games have the characteristics of:
  - Real-time decision making
  - Extremely large branching factor for search-based approaches
  - Partially-observable

# Reinforcement Learning



- At each step  $t$  the agent:
  - Executes action  $A_t$
  - Receives observation  $O_t$
  - Receives scalar reward  $R_t$
- The environment:
  - Receives action  $A_t$
  - Emits observation  $O_{t+1}$
  - Emits scalar reward  $R_{t+1}$
- $t$  increments at env. step

# Problem Statement

- **Problem:** how do we design observation and action representation of RTS games for DRL (Deep Reinforcement Learning)?
  - Observations
    - Raw pixels?
    - Discretized states (A binary array of attributes “haveBarracks”, “haveRanged”, “isAttacking”)?
  - Actions
    - Mouse clicks and key strokes?
    - For each unit, give a command (soilder1: attack, soilder2: noop)?

# Problem Statement


- **Problem:** how do we design observation and action representation of RTS games for DRL (Deep Reinforcement Learning)?

## This paper:

- We compared two intuitive representations:
  - a *global* representation where the agent *sees* everything and *controls* an interested unit at each timestep.
  - a *local* representation where the agent becomes an unit at each timestep, *sees* everything that is some distance away from it and *controls* itself.

$\mu RTS$

- Open-source <https://github.com/santiontanon/microrts>
- Minimalistic
- Deterministic
- Fully or partially-observable options
- Available AI techniques: UCT, NaiveMCTS, Portfolio greedy search, AHTN, ...





# Existing Work

RL techniques in RTS games:

- Alisp to specify a list of commands, and Q-learning to tune the parameters [Marthi et al., 2005] in Wargus.
- Q-learning to learn a policy for each class of units (peasants, knights, barracks, etc) [Marthi et al., 2005] in Wargus.
- Option framework and heuristic algorithms to simplify the game space [Tavares and Chaimowicz, 2018].

# Existing Work


Recent development utilizing neural networks with RL:

- Provide feature maps and low-level actions to train agents.  
[Vinyals et al., 2017]
- High-level actions and curriculum.  
training[Tian et al., 2017, Sun et al., 2018, Lee et al., 2018]
- HRL (Hierarchical Reinforcement Learning) [Liu et al., 2019].
- Multi-Agent representation [Samvelyan et al., 2019]

- 1 Introduction
  - Background
  - Problem Statement
  - Existing Work
- 2 Observation and Action Representations
  - Global Representation
  - Local Representation
- 3 Advantage Actor Critic
  - A2C
  - Specifics
- 4 Evaluation
  - Experimental Setup
  - Evaluation
- 5 Conclusions


# Global Representation

- Let the observation represents the whole game state



# Global Representation

- We have 5 feature maps: hit points, resources, ownership, unit type, and unit actions



$$\begin{bmatrix} 1 & 1 & 0 & 0 \\ 1 & 10 & 0 & 0 \\ 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 10 \end{bmatrix}$$

Hit points feature map




$$\begin{bmatrix} 0 & 1 & 0 & 0 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 & 0 & 0 & 0 \\ 1 & 0 & 0 & 0 & 0 & 0 & 0 \\ 1 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 1 \\ \vdots & \vdots & \vdots & \vdots & \vdots & \vdots & \vdots \\ 0 & 0 & 0 & 0 & 0 & 0 & 1 \end{bmatrix}$$

One-hot-encoded array


# Global Representation

- Let the RL agent choose which unit to issue actions to, and which actions to execute.




# Local Representation

- Let the observation to be represented from the point of view of an individual unit. Introduce the parameter of window size  $w$  that specifies the observable distance away from the selected unit.



# Local Representation

- Let the RL agent picks actions for each unit independently.





## Other Details and Insights

- If the action produced is invalid, then it will be replaced by an NOOP action.
- The agent can only issue one action to one unit at each timestep.
- The local representation rotates through units and asks the RL agent to produce actions for the selected unit only.

- 1 Introduction
  - Background
  - Problem Statement
  - Existing Work
- 2 Observation and Action Representations
  - Global Representation
  - Local Representation
- 3 **Advantage Actor Critic**
  - A2C
  - **Specifics**
- 4 Evaluation
  - Experimental Setup
  - Evaluation
- 5 Conclusions

# A2C

---

**Algorithm 1** Advantage Actor Critic

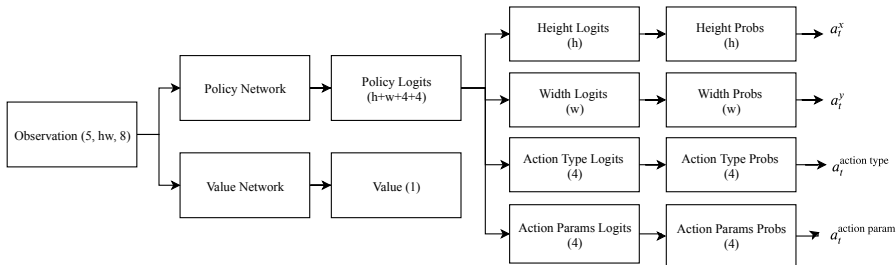
---

```
1: Initialize policy function  $\pi$  with random weights  $\theta$ 
2: Initialize value function  $v$  with random weights  $\theta'$ 
3: for episode = 1,  $M$  do
4:   Reset the game and get  $s_1$ 
5:   for  $t = 1, T$  do
6:     Sample action  $a_t$  from  $\pi(s_t)$ 
7:     Execute action  $a_t$  and record reward  $r_t$  and state  $s_{t+1}$ 
8:     If  $s_{t+1}$  is terminal, break
9:   end for
10:  for  $t = 1, T$  do
11:    Calculate the value  $v_t$ 
12:    Calculate the advantage  $A = \sum_{i=1}^T \gamma^i r_t - v_t$ 
13:     $\theta = \theta + \alpha \nabla_{\theta} A \log \pi(a_t | s_t) + \pi_{\theta}(a_t | s_t) \log \pi_{\theta}(a_t | s_t)$ 
14:     $\theta' = \theta' + \beta \nabla_{\theta'} A^2$ 
15:  end for
16: end for
```

---

# Specifics

How exactly do we input  $s_t$  to A2C and generate action  $a_t$ ?



# Specifics

How exactly do we calculate  $\log \pi_{\theta}(a_t|s_t)$  and  $\pi_{\theta}(a_t|s_t) \log \pi_{\theta}(a_t|s_t)$ ?

$$\log \pi_{\theta}(a_t|s_t) = \log \pi_{\theta}(a_t^x|s_t)$$

$$+ \log \pi_{\theta}(a_t^y|s_t)$$

$$+ \log \pi_{\theta}(a_t^{\text{action type}}|s_t)$$

$$+ \log \pi_{\theta}(a_t^{\text{action param}}|s_t)$$

$$\pi_{\theta}(a_t|s_t) \log \pi_{\theta}(a_t|s_t) = \pi_{\theta}(a_t^x|s_t) \log \pi_{\theta}(a_t^x|s_t)$$

$$+ \pi_{\theta}(a_t^y|s_t) \log \pi_{\theta}(a_t^y|s_t)$$




$$+ \pi_{\theta}(a_t^{\text{action type}}|s_t) \log \pi_{\theta}(a_t^{\text{action type}}|s_t)$$

$$+ \pi_{\theta}(a_t^{\text{action param}}|s_t) \log \pi_{\theta}(a_t^{\text{action param}}|s_t)$$

- 1 Introduction
  - Background
  - Problem Statement
  - Existing Work
- 2 Observation and Action Representations
  - Global Representation
  - Local Representation
- 3 Advantage Actor Critic
  - A2C
  - Specifics
- 4 **Evaluation**
  - Experimental Setup
  - Evaluation
- 5 Conclusions


# Experimental Setup

- Task: harvest resources with different map sizes of  $4 \times 4$ ,  $6 \times 6$ , and  $8 \times 8$ .
- It takes 10 timesteps to execute actions (move, harvest, return)
- When the agent harvested or returned the resources, it gets a reward of 10. Otherwise the reward is 0.



## Episode Rewards

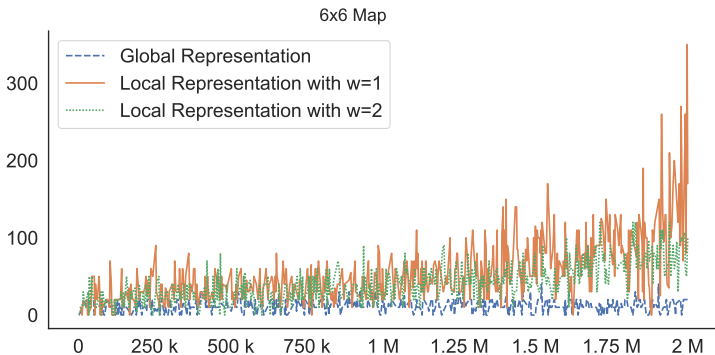
Episode rewards (y axis) as a function of training time steps (x-axis) for the 3 map sizes. (Higher is better)






# Episode Rewards

Episode rewards (y axis) as a function of training time steps (x-axis) for the 3 map sizes. (Higher is better)



# Episode Rewards

Episode rewards (y axis) as a function of training time steps (x-axis) for the 3 map sizes. (Higher is better)



# Metrics Evaluation

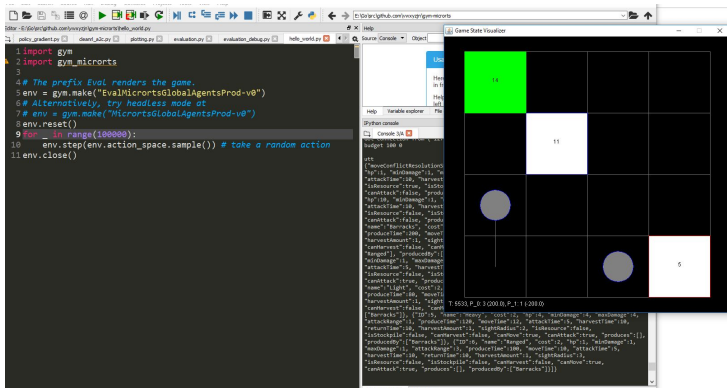
	map	$t_{\text{first harvest}}$	$t_{\text{first return}}$	$r$
RandomAI	$4 \times 4$	51.33	142.67	11.87
Global	$4 \times 4$	99.00	167.73	13.13
Local ( $w = 1$ )	$4 \times 4$	29.87	172.47	67.20
Local ( $w = 2$ )	$4 \times 4$	45.00	73.73	33.40
RandomAI	$6 \times 6$	421.33	797.33	2.00
Global	$6 \times 6$	533.33	1931.20	0.07
Local ( $w = 1$ )	$6 \times 6$	59.20	567.40	3.53
Local ( $w = 2$ )	$6 \times 6$	62.33	408.73	3.93
RandomAI	$8 \times 8$	878.67	1480.67	0.87
Global	$8 \times 8$	1464.53	-	0.00
Local ( $w = 1$ )	$8 \times 8$	167.20	1844.20	0.20
Local ( $w = 2$ )	$8 \times 8$	89.87	-	0.00

# Visualization of Agents

<https://youtu.be/--BoB0wnF0s>

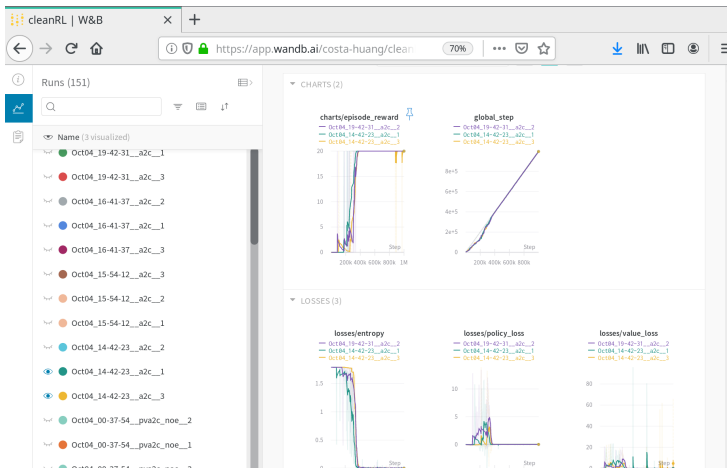
# Gym environment

Our OpenAI Gym library is available at  
<https://github.com/vwxyzjn/gym-microrts>



# DRL Library (WIP)

Our RL library is available at <https://github.com/vwxyzjn/cleanrl>



# Conclusions

- We compared two intuitive representations
  - a *global* representation that feeds the RL agent the entire game state, and require the agent to learn to locate the unit and control it.
  - a *local* representation that feeds the RL agent the localized game state that is some distance away from the selected unit, and require the agent to learn to control it
- We show that local representation generally outperforms but the training of agents becomes more difficult in larger maps, where the exploration and sparse rewards become a huge problem.

## Future Work

- Frame skipping
- Parallel processing
- Larger maps
- DQN-based algorithms
- Partial observability and LSTM
- Self-play



## Our Code

- <https://github.com/vwxyzjn/gym-microrts>
- <https://github.com/vwxyzjn/cleanrl>
- <https://github.com/vwxyzjn/microrts>

Thank you. Questions?

## References I



Lee, D., Tang, H., Zhang, J. O., Xu, H., Darrell, T., and Abbeel, P. (2018).

Modular architecture for starcraft ii with deep reinforcement learning.  
In *Fourteenth Artificial Intelligence and Interactive Digital Entertainment Conference*.



Liu, R.-Z., Guo, H., Ji, X., Yu, Y., Xiao, Z., Wu, Y., Pang, Z.-J., and Lu, T. (2019).

Efficient reinforcement learning with a mind-game for full-length starcraft ii.

*arXiv preprint arXiv:1903.00715*.

## References II



Marthi, B., Russell, S. J., Latham, D., and Guestrin, C. (2005).  
Concurrent hierarchical reinforcement learning.  
In *IJCAI*, pages 779–785.



Samvelyan, M., Rashid, T., de Witt, C. S., Farquhar, G., Nardelli, N.,  
Rudner, T. G., Hung, C.-M., Torr, P. H., Foerster, J., and Whiteson,  
S. (2019).  
The starcraft multi-agent challenge.  
*arXiv preprint arXiv:1902.04043*.

## References III



Sun, P., Sun, X., Han, L., Xiong, J., Wang, Q., Li, B., Zheng, Y., Liu, J., Liu, Y., Liu, H., et al. (2018).

Tstarbots: Defeating the cheating level builtin ai in starcraft ii in the full game.

*arXiv preprint arXiv:1809.07193.*





Tavares, A. R. and Chaimowicz, L. (2018).

Tabular reinforcement learning in real-time strategy games via options.

*2018 IEEE Conference on Computational Intelligence and Games (CIG), pages 1–8.*

## References IV

-  Tian, Y., Gong, Q., Shang, W., Wu, Y., and Zitnick, C. L. (2017). Elf: An extensive, lightweight and flexible research platform for real-time strategy games. In *Advances in Neural Information Processing Systems*, pages 2659–2669.
-  Vinyals, O., Ewalds, T., Bartunov, S., Georgiev, P., Vezhnevets, A. S., Yeo, M., Makhzani, A., Küttler, H., Agapiou, J., Schrittwieser, J., et al. (2017). Starcraft ii: A new challenge for reinforcement learning. *arXiv preprint arXiv:1708.04782*.